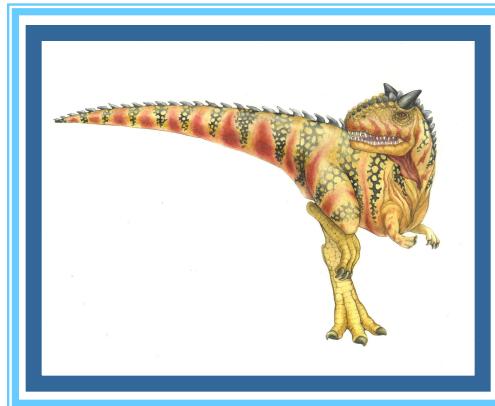
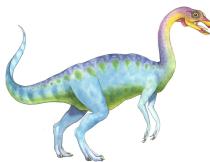


Chapter 8: Main Memory





Recap

- Background
- Address Binding
 - Base and Limit registers
 - Hardware Address Protection
- Logical and Physical addresses
- Dynamic Relocation (MMU)
- Dynamic Linking





Objectives

- Swapping
- Contiguous Memory Allocation
 - Fixed size partitioning
 - Variable size partitioning
- Dynamic Storage-Allocation Problem
 - First Fit
 - Best Fit
 - Worst Fit
- Fragmentation
 - Internal Fragmentation
 - External Fragmentation
- Compaction





Swapping

- A process can be **swapped** temporarily out of memory to a backing store, and then brought back into memory for continued execution
 - Total physical memory space of processes can exceed physical memory
- **Backing store** – fast disk large enough to accommodate copies of all memory images for all users; must provide direct access to these memory images
- **Roll out, roll in** – swapping variant used for priority-based scheduling algorithms; lower-priority process is swapped out so higher-priority process can be loaded and executed
- Major part of swap time is transfer time; total transfer time is directly proportional to the amount of memory swapped
- System maintains a **ready queue** of ready-to-run processes which have memory images on disk





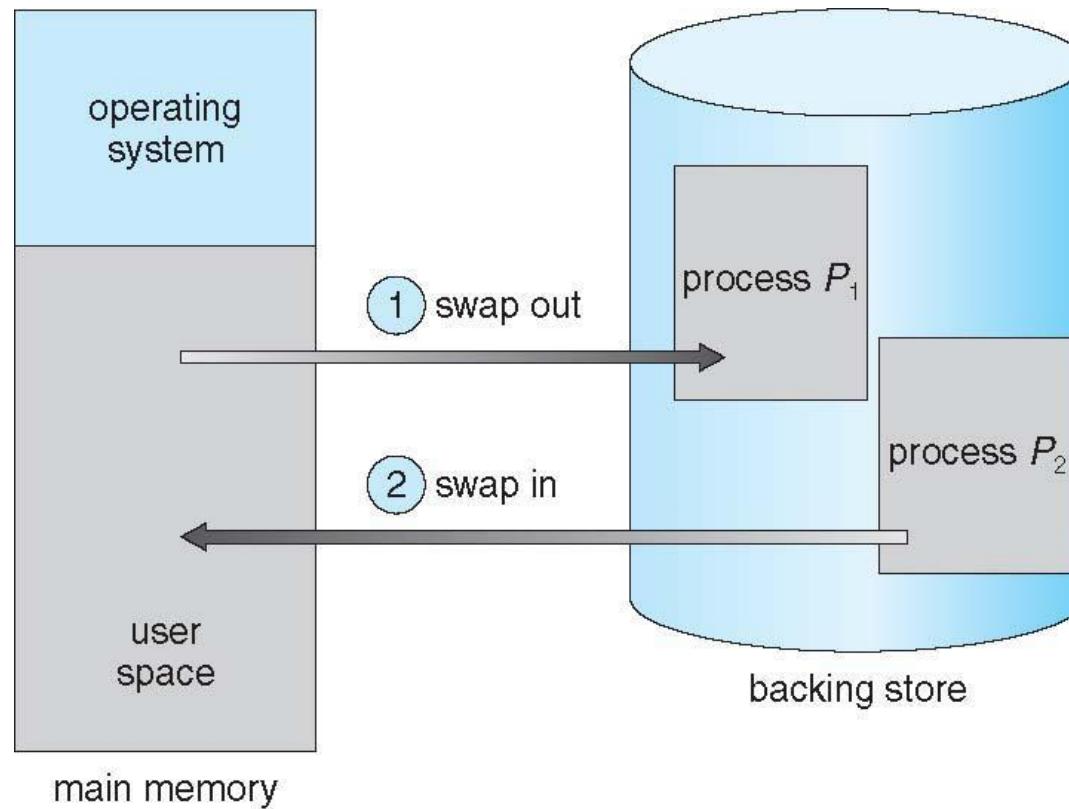
Swapping (Cont.)

- Does the swapped-out process need to swap back into same physical addresses?
- Depends on address binding method
 - Plus consider pending I/O to / from process memory space
- Modified versions of swapping are found on many systems (i.e., UNIX, Linux, and Windows)
 - Swapping normally disabled
 - Started if more than threshold amount of memory allocated
 - Disabled again once memory demand reduced below threshold



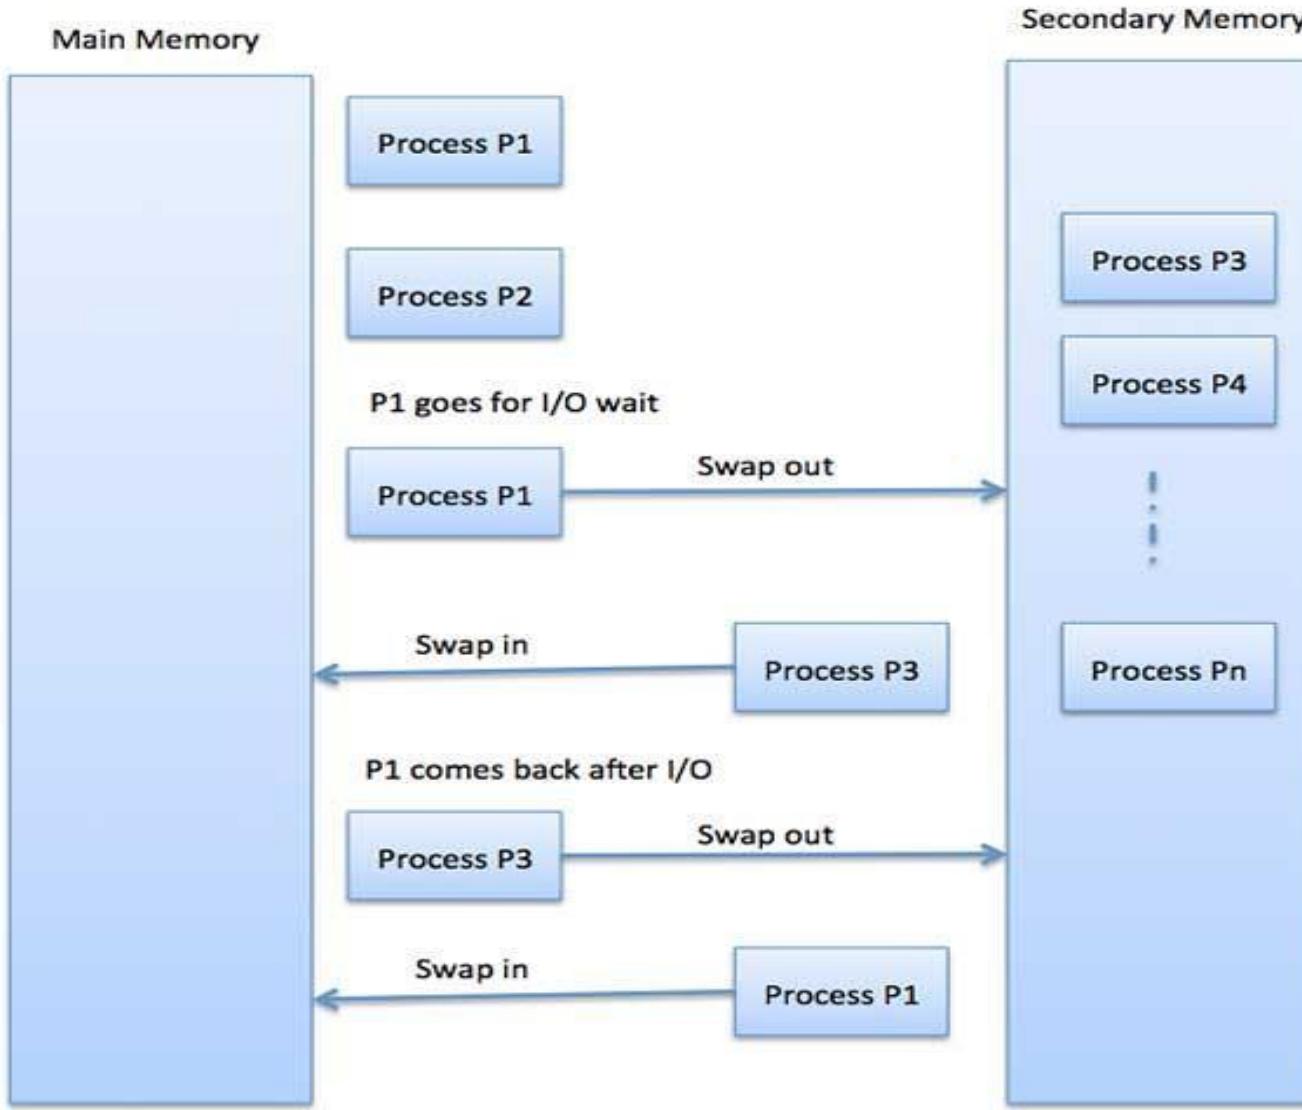


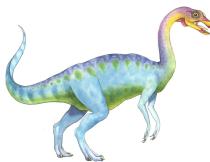
Schematic View of Swapping





Schematic View of Swapping





Context Switch Time including Swapping

- If next processes to be put on CPU is not in memory, need to swap out a process and swap in target process
- Context switch time can then be very high
- 100MB process swapping to hard disk with transfer rate of 50MB/sec
 - Swap out time of 2000 ms
 - Plus swap in of same sized process
 - Total context switch swapping component time of 4000ms (4 seconds)
- Can reduce if reduce size of memory swapped – by knowing how much memory really being used
 - System calls to inform OS of memory use via `request_memory()` and `release_memory()`

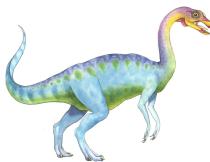




Context Switch Time and Swapping (Cont.)

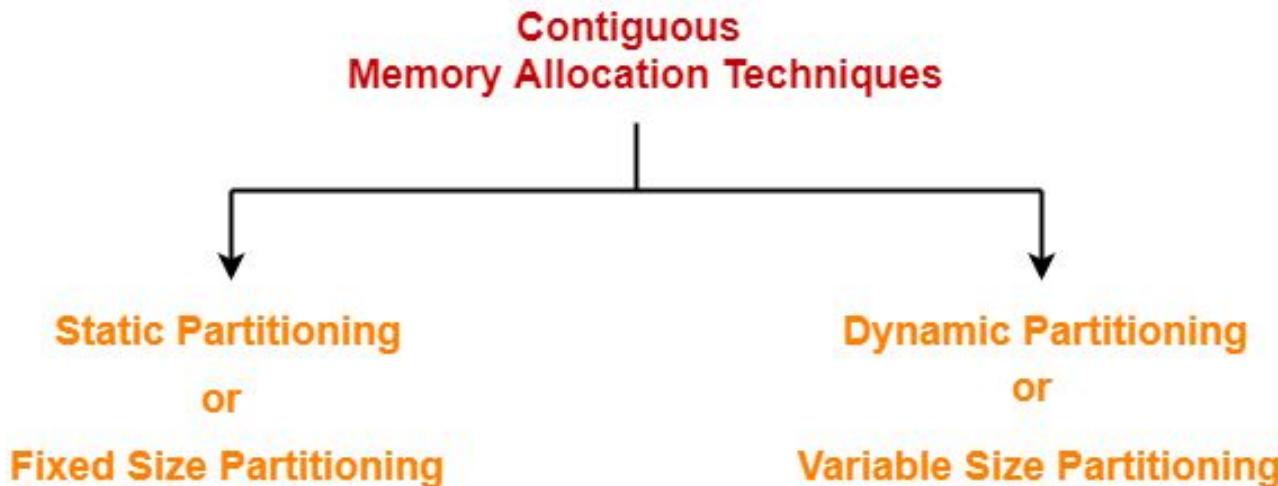
- Other constraints as well on swapping
 - Pending I/O – can't swap out as I/O would occur to wrong process
 - Or always transfer I/O to kernel space, then to I/O device
 - 4 Known as **double buffering**, adds overhead
- Standard swapping not used in modern operating systems
 - But modified version common
 - 4 Swap only when free memory extremely low





Contiguous Allocation

- Main memory must support both OS and user processes
- Limited resource, must allocate efficiently
- Contiguous allocation is one early method
- Main memory usually divided into two **partitions**:
 - Resident operating system, usually held in low memory with interrupt vector
 - User processes then held in high memory
 - Each process contained in single contiguous section of memory





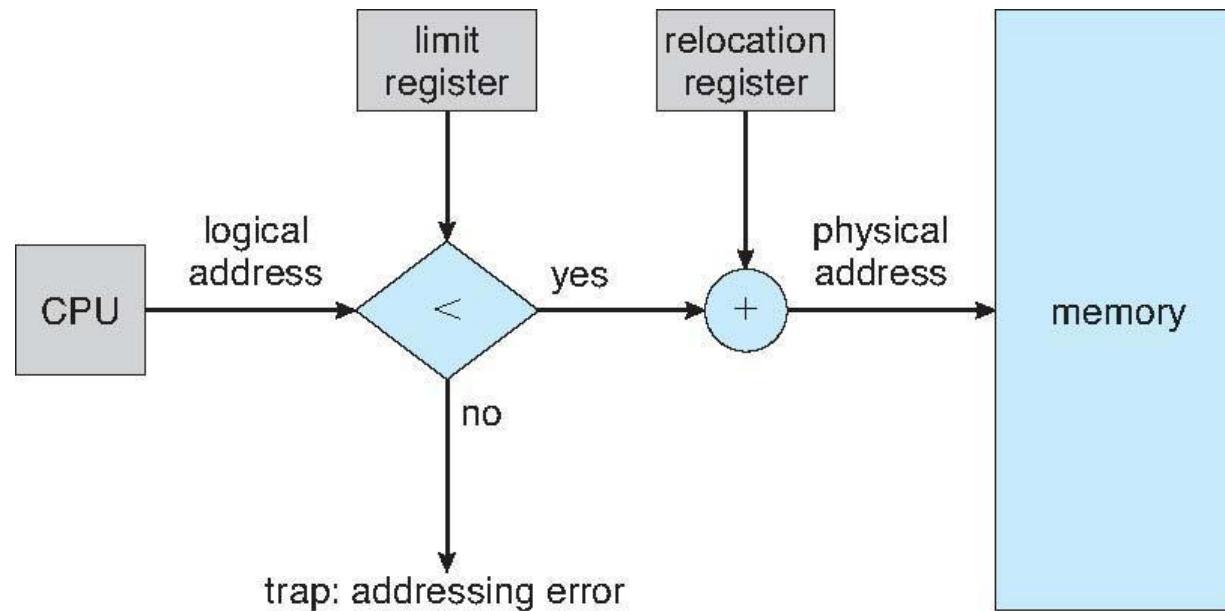
Contiguous Allocation (Cont.)

- Relocation registers used to protect user processes from each other, and from changing operating-system code and data
 - Base register contains value of smallest physical address
 - Limit register contains range of logical addresses – each logical address must be less than the limit register
 - MMU maps logical address *dynamically*
 - Can then allow actions such as kernel code being **transient** and kernel changing size





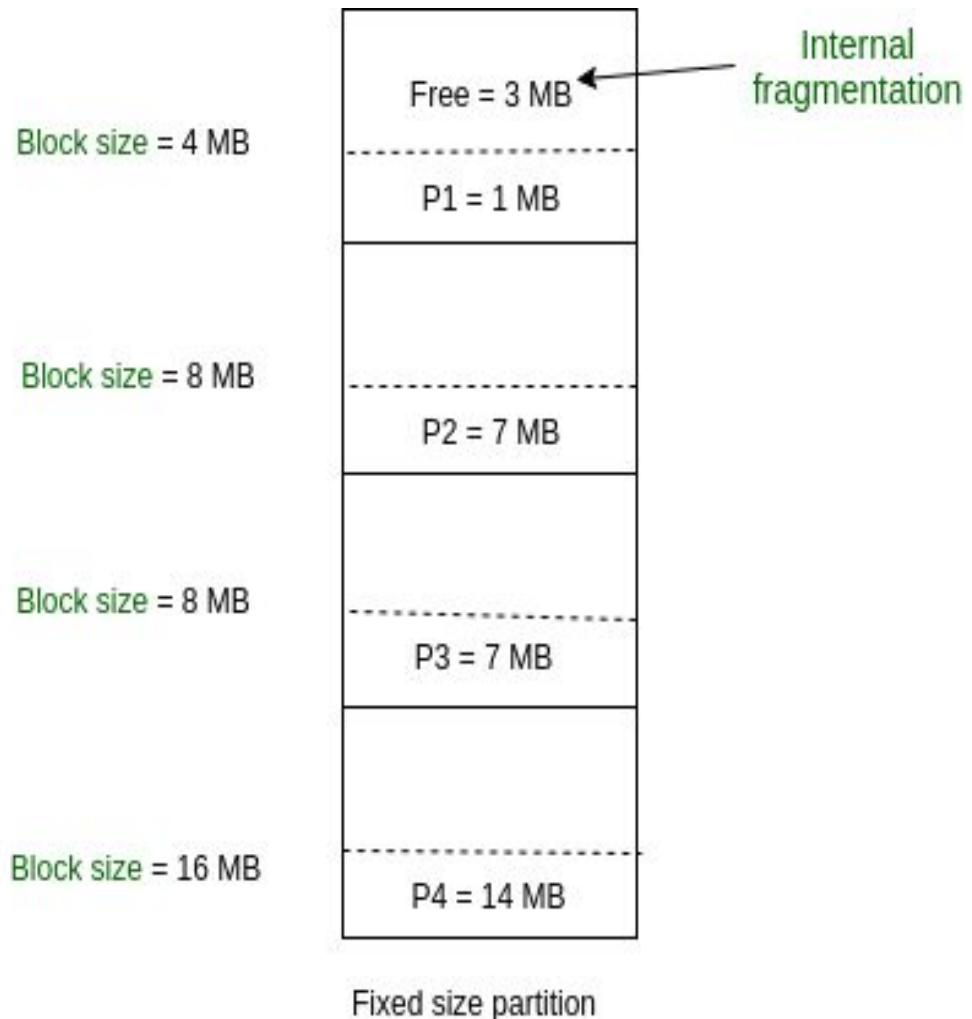
Hardware Support for Relocation and Limit Registers

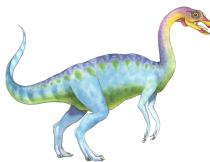




Fixed Size Partitioning

- This is the oldest and simplest technique used to put more than one processes in the main memory. In this partitioning, number of partitions (non-overlapping) in RAM are **fixed but size** of each partition may or **may not be same**.
- As it is **contiguous** allocation, hence no spanning is allowed. Here partition are made before execution or during system configure.+/

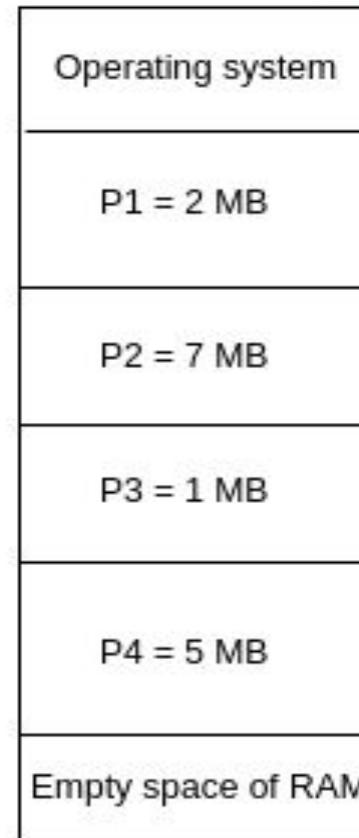




Variable Size Partitioning

- In contrast with fixed partitioning, partitions are not made before the execution or during system config. Various **features** associated with variable Partitioning-
- Initially RAM is empty, and partitions are made during the run-time according to process's need instead of partitioning during system config.
- The size of partition will be equal to incoming process.
- The partition size varies according to the need of the process so that the internal fragmentation can be avoided to ensure efficient utilization of RAM.
- Number of partitions in RAM is not fixed and depends on the number of incoming process and Main Memory's size.

Dynamic partitioning



Block size = 2 MB

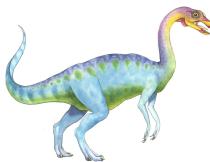
Block size = 7 MB

Block size = 1 MB

Block size = 5 MB

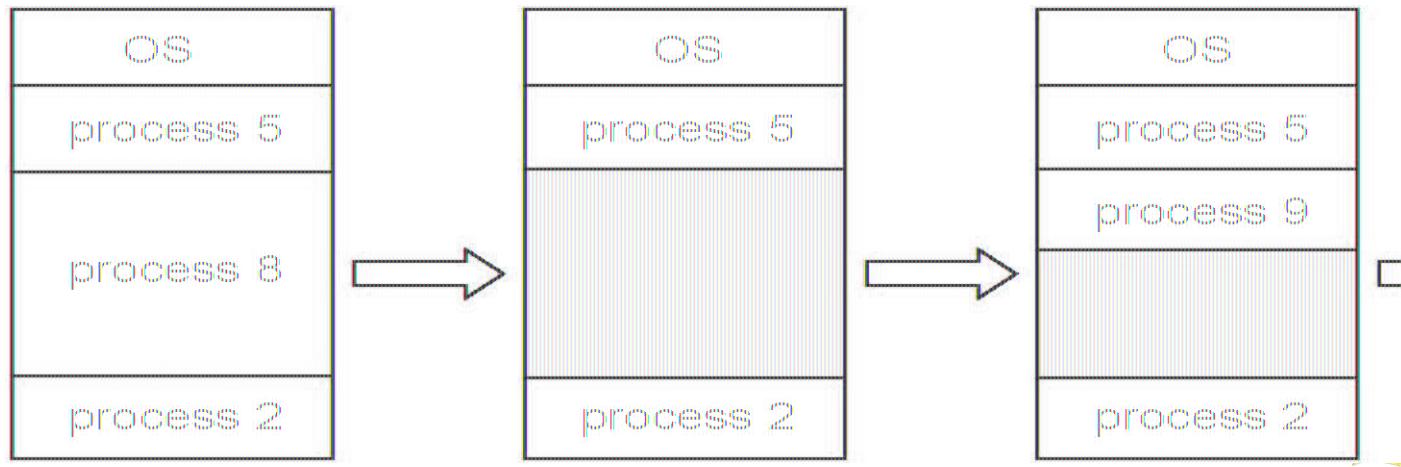
Partition size = process size
So, no internal Fragmentation





Multiple-partition allocation

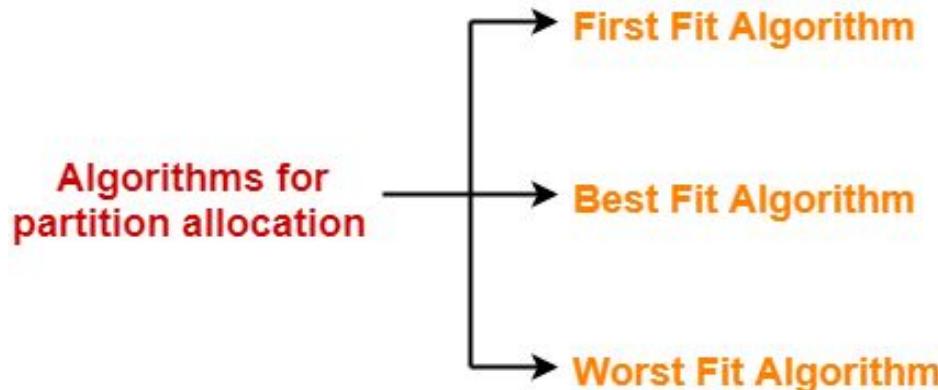
- Multiple-partition allocation
 - Degree of multiprogramming limited by number of partitions
 - **Variable-partition** sizes for efficiency (sized to a given process' needs)
 - **Hole** – block of available memory; holes of various size are scattered throughout memory
 - When a process arrives, it is allocated memory from a hole large enough to accommodate it
 - Process exiting frees its partition, adjacent free partitions combined
 - Operating system maintains information about:
 - a) allocated partitions
 - b) free partitions (hole)





Dynamic Storage-Allocation Problem

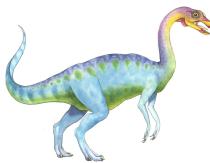
How to satisfy a request of size n from a list of free holes?



- **First-fit**: Allocate the *first* hole that is big enough
- **Best-fit**: Allocate the *smallest* hole that is big enough; must search entire list, unless ordered by size
 - Produces the smallest leftover hole
- **Worst-fit**: Allocate the *largest* hole; must also search entire list
 - Produces the largest leftover hole

First-fit and best-fit better than worst-fit in terms of speed and storage utilization





First Fit

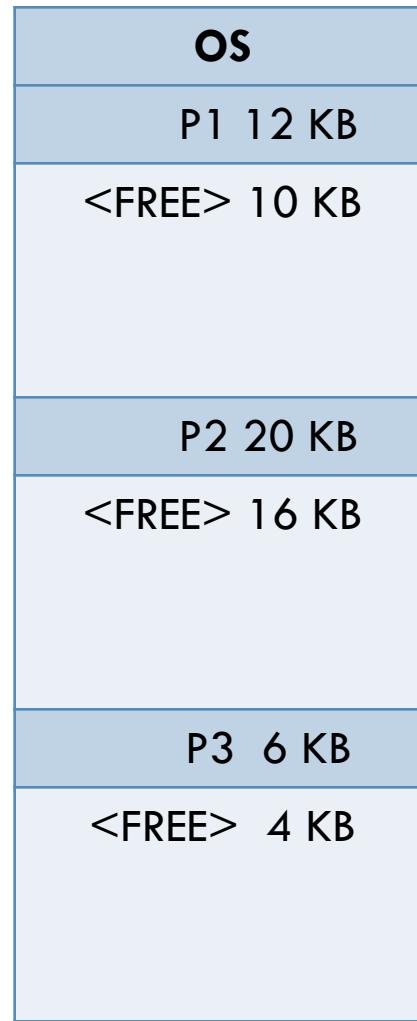
- First Fit : Allocate the first free block that is large enough for the new process.
- This is a fast algorithm.

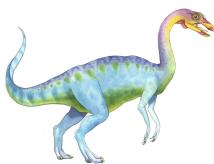




First Fit Example

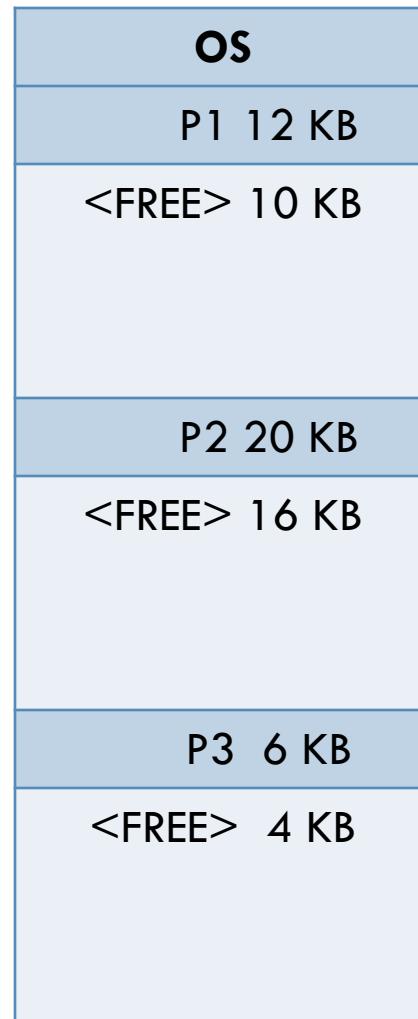
Initial memory
mapping

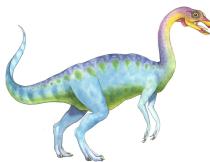




First Fit Example (Cont.)

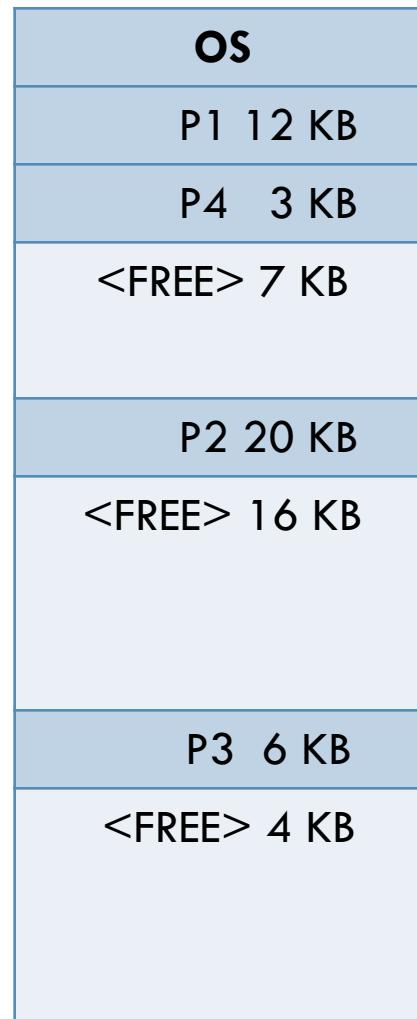
P4 of 3KB
arrives





First Fit Example (Cont.)

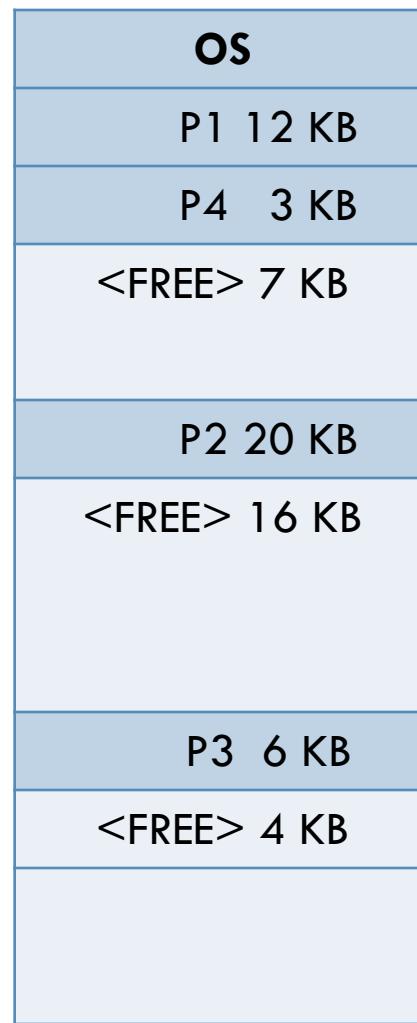
P4 of 3KB
loaded here by
FIRST FIT





First Fit Example (Cont.)

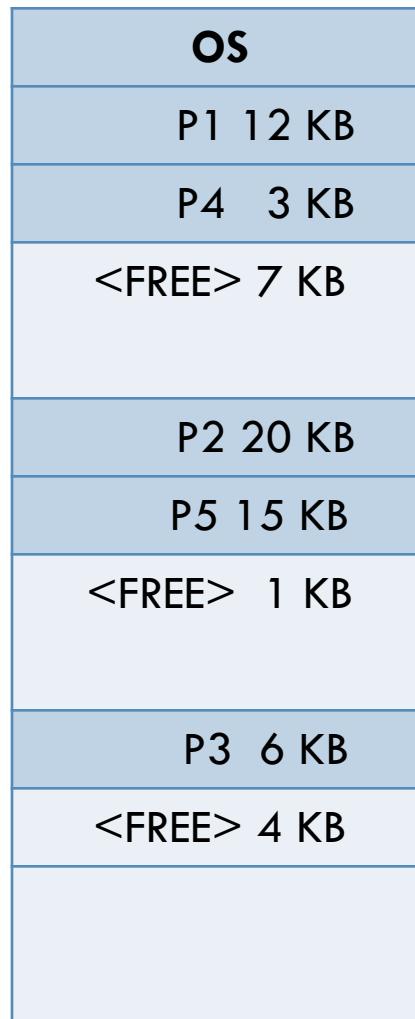
P5 of 15KB
arrives

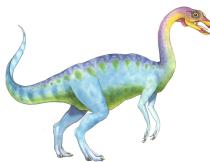




First Fit Example (Cont.)

P5 of 15 KB
loaded here by
FIRST FIT

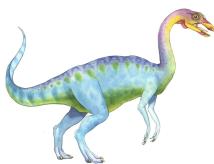




Best fit

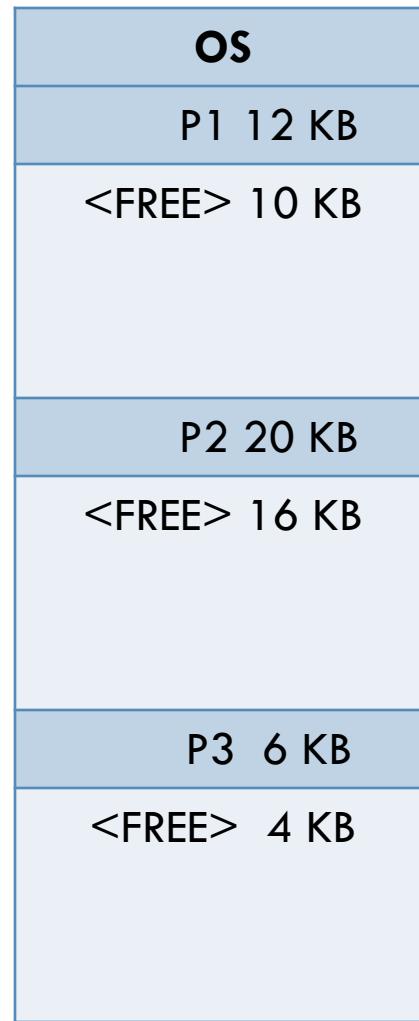
- Best Fit : Allocate the smallest block among those that are large enough for the new process.
- In this method, the OS has to search the entire list, or it can keep it sorted and stop when it hits an entry which has a size larger than the size of new process.
- This algorithm produces the smallest left over block.
- However, it requires more time for searching all the list or sorting it
- If sorting is used, merging the area released when a process terminates to neighboring free blocks, becomes complicated.





Best Fit Example

Initial memory mapping





Best Fit Example (Cont.)

P4 of 3KB
arrives

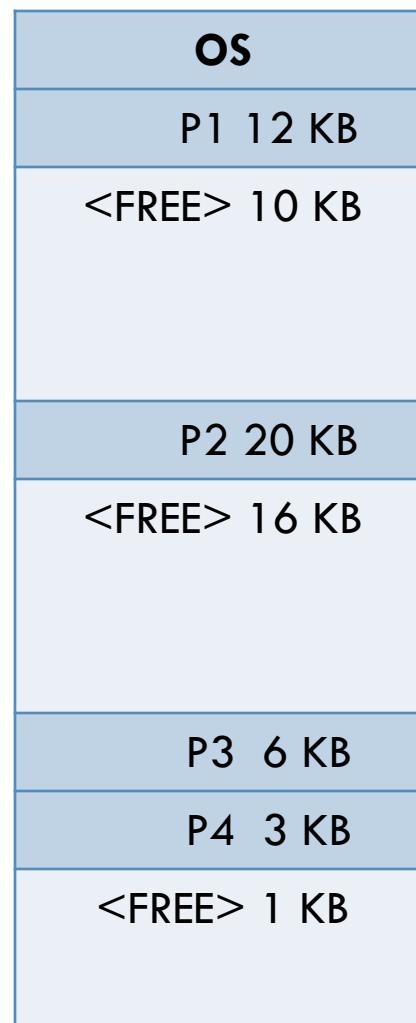
OS
P1 12 KB
<FREE> 10 KB
P2 20 KB
<FREE> 16 KB
P3 6 KB
<FREE> 4 KB





Best Fit Example (Cont.)

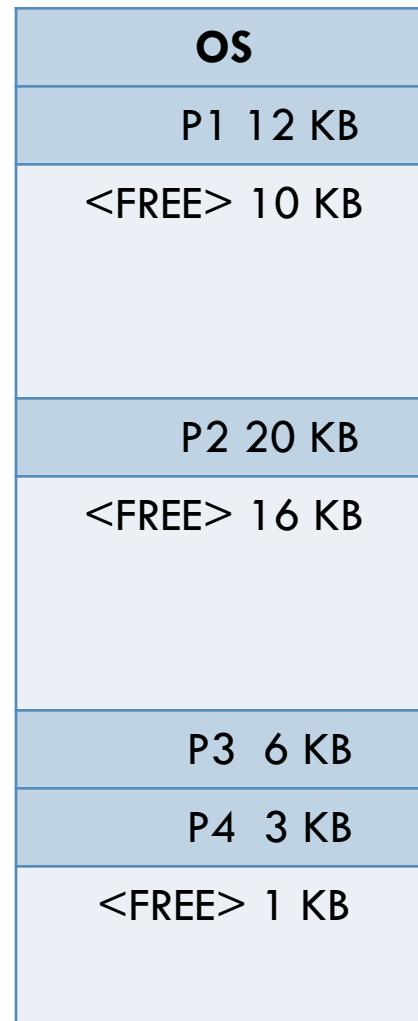
P4 of 3KB
loaded here by
BEST FIT





Best Fit Example (Cont.)

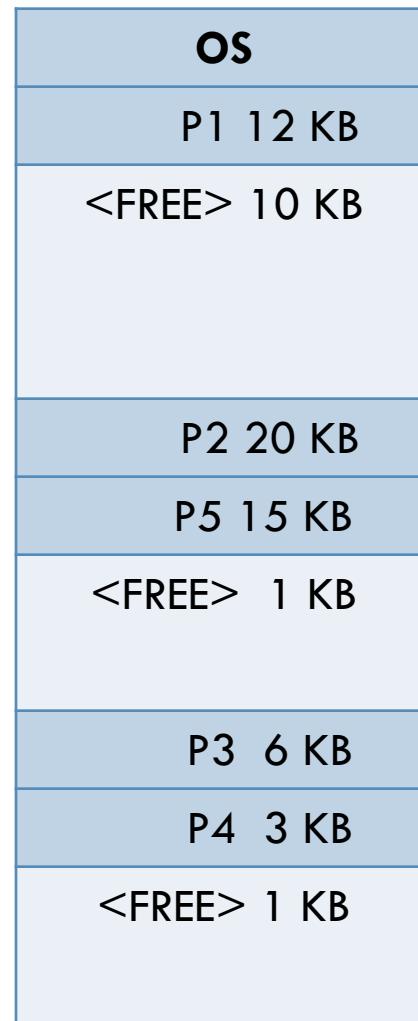
P5 of 15KB
arrives

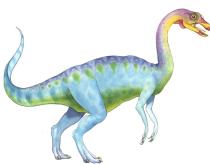




Best Fit Example (Cont.)

P5 of 15 KB
loaded here by
BEST FIT

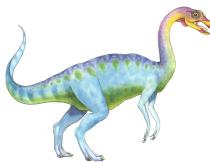




Worst Fit

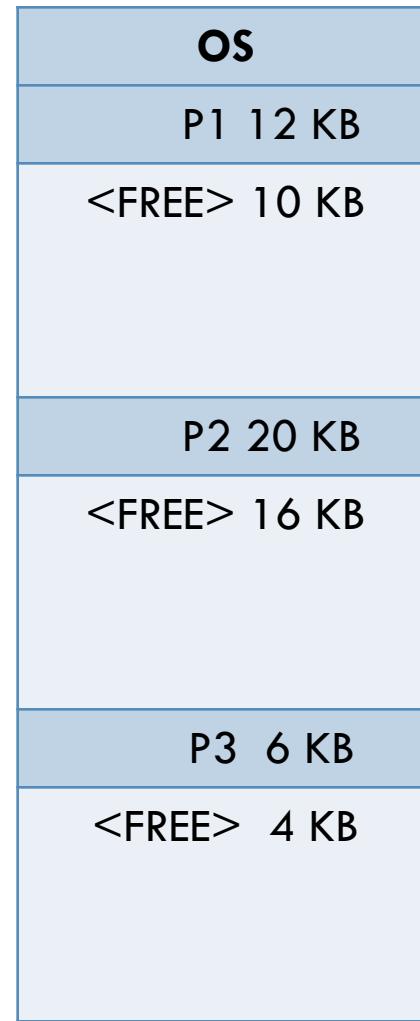
- Worst Fit : Allocate the largest block among those that are large enough for the new process.
- Again a search of the entire list or sorting it is needed.
- This algorithm produces the largest over block.





Worst Fit Example

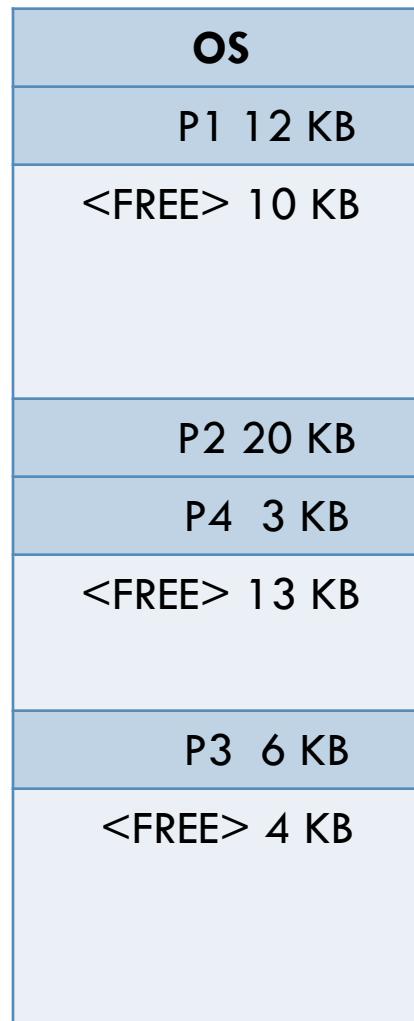
Initial memory mapping





Worst Fit Example (Cont.)

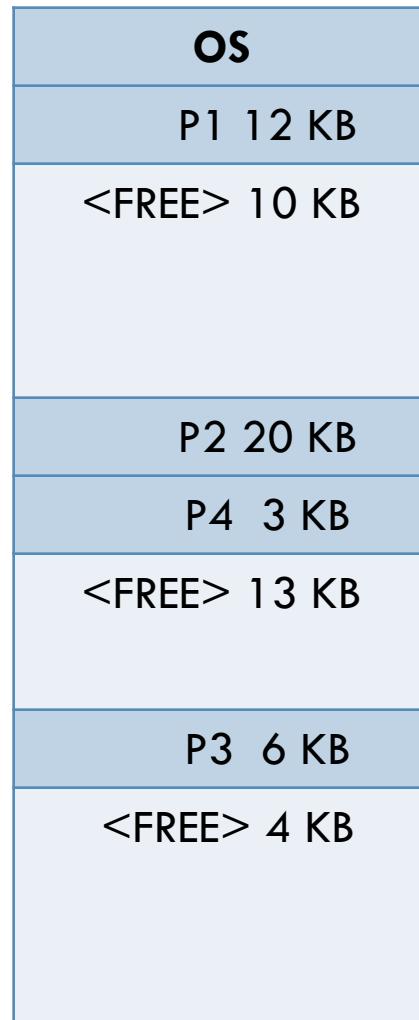
P4 of 3KB
Loaded here by
WORST FIT





Worst Fit Example (Cont.)

No place to load
P5 of 15K





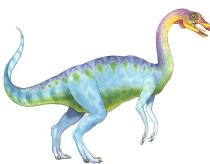
Worst Fit Example (Cont.)

No place to load
P5 of 15K

Compaction is
needed !!

OS
P1 12 KB
<FREE> 10 KB
P2 20 KB
P4 3 KB
<FREE> 13 KB
P3 6 KB
<FREE> 4 KB

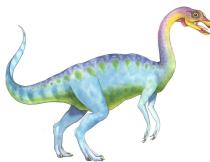




Fragmentation

- **External Fragmentation** – total memory space exists to satisfy a request, but it is not contiguous
- **Internal Fragmentation** – allocated memory may be slightly larger than requested memory; this size difference is memory internal to a partition, but not being used





Fragmentation



If a whole partition is currently not being used, then it is called an **external fragmentation**.

If a partition is being used by a process requiring some memory smaller than the partition size, then it is called an **internal fragmentation**.





Fragmentation (Cont.)

- Reduce external fragmentation by **compaction**
 - Shuffle memory contents to place all free memory together in one large block
 - Compaction is possible *only* if relocation is dynamic, and is done at execution time
 - I/O problem
 - 4 Latch job in memory while it is involved in I/O
 - 4 Do I/O only into OS buffers
- Now consider that backing store has same fragmentation problems



